Lecture 1: Intro to simulation

Warm-up question

Problem: 10 people are at a party, and all of them are wearing hats. They each place their hat in a pile; when they leave, they choose a hat at random. What is the probability at least one person selected the correct hat?

Question: Work with your neighbor to discuss the following question:

- Without calculating probabilities, how could you design an experiment to estimate this probability?

Designing an experiment

represent hats w/ #s 1–10

Step 1: need 10 hats! (each belongs to one person)

person 1 person 2 3 ... person 10

Hat 1 Hat 2 Hat 3 ... Hat 10

Step 2: randomly "shuffle" hats / sampling in R

randomly assign a hat to each person

10 4 3 9

Step 3: who got their original hat?

Step 4: repeat many times!

for loop!

Step 1: representing the hats

```
1 hats <- 1:10

2

3 hats
```

vector called "hats"

vector 1, 2, 3, ..., 10

store the right hand side in the left hand side

```
 [1]  1  2  3  4  5  6  7  8  9 10
```

```
1 hats[3]
```

```
[1] 3
```

third entry

- `hats` is a **vector**, containing the numbers 1 to 10
- entries in a vector are accessed by their index

Step 2: everyone draws a random hat

```
1 hats <- 1:10
2 randomized_hats <- sample(hats, size = 10, replace = FALSE)
3
4 hats
```

```
 [1]  1  2  3  4  5  6  7  8  9 10
```

```
1 randomized_hats
```

```
 [1]  6  7 10  5  2  1  3  8  9  4
```

- The `sample` function creates a random sample from a vector
- How many people selected their original hat?

2 people.

Step 3: check who got their original hat

```
hats <- 1:10
randomized_hats <- sample(hats, size = 10, replace = FALSE)
```

```
hats
```

```
 [1]  1  2  3  4  5  6  7  8  9 10
```

```
randomized_hats
```

```
 [1]  6  7 10  5  2  1  3  8  9  4
```

```
hats == randomized_hats
```

```
 [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE  TRUE  TRUE FALSE
```

== test for equality (done for each element)

```
# TRUE is 1, FALSE is 0
sum(hats == randomized_hats)
```

```
[1] 2
```

how many people got their original hat?

```
# did at least one person get their hat?
sum(hats == randomized_hats) > 0
```

```
[1] TRUE
```

(at least one person got their original hat)

Code so far

```
1 hats <- 1:10
2 randomized_hats <- sample(hats, size = 10, replace = FALSE)
```

```
1 sum(hats == randomized_hats) > 0
```

```
[1] TRUE
```

- Is this a good estimate of the probability?

No! Need to repeat many times

Step 4: iteration

A `for` loop repeats code many times:

```
1 nsim <- 10000 # number of simulations
2 for(i in 1:nsim){
3
4       "repeat the following code for i=1,...,nsim"
5 }
```

i=1

(run code)

i=2

(run code)

i=3

(run code)

⋮

Step 4: iteration

results <- rep(NA, nsim)
NA NA NA ..... NA (10000 times)

A `for` loop repeats code many times:

```
1  nsim <- 10000 # number of simulations
2  hats <- 1:10
3  results <- rep(NA, nsim) # vector to store results
4
5  for(i in 1:nsim){
6    randomized_hats <- sample(hats, size = 10, replace = FALSE)
7    results[i] <- sum(hats == randomized_hats) > 0
8  }
9
10 head(results)
```

(hats 1, ..., 10)

"repeat NA nsim times"

results vector (store what happens)

store results

check whether each person got their original hat

```
[1] FALSE  TRUE  TRUE FALSE  TRUE  TRUE
```

prob. = proportion of TRUE in results vector

NA NA NA ..... NA

i = 1:
(check)
FALSE NA NA NA .....

i = 2:
(check)
FALSE TRUE NA NA .....

i = 10000

Step 4: iteration

A `for` loop repeats code many times:

```
 1  nsim <- 10000 # number of simulations
 2  hats <- 1:10
 3  results <- rep(NA, nsim) # vector to store results
 4
 5  for(i in 1:nsim){
 6    randomized_hats <- sample(hats, size = 10, replace = FALSE)
 7    results[i] <- sum(hats == randomized_hats) > 0
 8  }
 9
10  mean(results)
```

magic numbers

```
[1] 0.6357
```

P(at least one person receives original hat) ≈ 0.636

- What if I wanted to repeat the simulation, with a different number of people?

Removing magic numbers

Without magic numbers:

```
 1 nsim <- 10000 # number of simulations
 2 M <- 10 # number of people
 3 hats <- 1:M
 4 results <- rep(NA, nsim) # vector to store results
 5
 6 for(i in 1:nsim){
 7   randomized_hats <- sample(hats,
 8                             size = M,
 9                             replace = FALSE)
10   results[i] <- sum(hats ==
11                       randomized_hats) > 0
12 }
13
14 mean(results)
```

```
[1] 0.6269
```

- Why did I get different results?

Final code

seed

setting a seed: each time I run a process w/ "random" functions, I get the same result

```
1  set.seed(3) # set a seed for reproducibility
2
3  M <- 10 # number of people at the party
4  hats <- 1:M # numbered hats
5  nsim <- 10000 # number of simulations
6  results <- rep(NA, nsim) # vector to store the results
7
8  for(i in 1:nsim){
9    # hats are randomly assigned to each person
10   randomized_hats <- sample(hats, M, replace = F)
11
12   # did at least one person get their hat back?
13   results[i] <- sum(randomized_hats == hats) > 0
14 }
15
16 mean(results)
```

Summary of coding practices

- avoid magic numbers
- set a seed for reproducibility
- use meaningful names
- add comments

Class activity

Work with a neighbor on the class activity (link below and on the course website):

https://sta279-s24.github.io/class_activities/ca_lecture_1.html