

Lecture 9: Rectangular data in R

Content so far

- Simulation
- Iteration
- Vectors and lists
- Functions

What's missing: actual data sets!

Learning goals

- Review/refresh data manipulation from STA 112
- Explore different data objects in R and Python
- Work with more challenging data, requiring more difficult manipulation
- Combine information from multiple datasets
- Learn tools for different data types (strings, factors, dates and times)

Next steps

- Today: overview of matrices and data frames in R; basic data wrangling in R
- Next week: Introduction to Python
- After that: data wrangling with both R and Python

Rectangular data

```
1 library(dplyr)
2 starwars
```

```
# A tibble: 87 × 14
```

	name	height	mass	hair_color	skin_color	eye_color	birth_year	sex
gender	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>	<dbl>	
	<chr> <chr>							
	1 Luke Sk...	172	77	blond	fair	blue	19	male
mascu...								
	2 C-3PO	167	75	<NA>	gold	yellow	112	none
mascu...								
	3 R2-D2	96	32	<NA>	white, bl...	red	33	none
mascu...								
	4 Darth V...	202	136	none	white	yellow	41.9	male
mascu...								
	5 Leia Or...	150	49	brown	light	brown	19	

Rectangular data in R

In R, there are two main ways of storing rectangular data:

- matrices
- data frames

Matrices

A *matrix* generalizes a vector to *two* dimensions:

```
1 x <- matrix(c(1, 2, 3, 4, 5, 6), nrow=2)
2 x
```

```
      [,1] [,2] [,3]
[1,]    1    3    5
[2,]    2    4    6
```

- Each row is a vector
- Each column is a vector

Indexing matrices

```
1 x <- matrix(c(1, 2, 3, 4, 5, 6), nrow=2)
2 x
```

	[,1]	[,2]	[,3]
[1,]	1	3	5
[2,]	2	4	6

```
1 x[1,]
```

```
[1] 1 3 5
```

```
1 x[,1]
```

```
[1] 1 2
```

```
1 x[1,2]
```

```
[1] 3
```

- Use single square brackets [] to index
- The first coordinate is the row, the second coordinate is the column

Uses and limitations of matrices

- Correspond to the matrices we know and love from linear algebra
- Usually the right way to store 2-d data for doing math (like matrix multiplication)
- Like vectors, contain only one type of data

```
1 x <- matrix(c(1, 2, 3, 'a', 5, 6), nrow=2)
2 x
```

```
      [,1] [,2] [,3]
[1,] "1"  "3"  "5"
[2,] "2"  "a"  "6"
```

Data frames

```
1 example_df <- data.frame(x = c(1, 2, 3),  
2                           y = c('a', 'b', 'c'))  
3 example_df
```

	x	y
1	1	a
2	2	b
3	3	c

Aside: what *are* data frames?

```
1 example_df <- data.frame(x = c(1, 2, 3),  
2                           y = c('a', 'b', 'c'))  
3 typeof(example_df)
```

```
[1] "list"
```

- Matrices are like a 2-d vector
- Data frames are a special type of list! With some requirements:
 - Each component is a vector
 - Each component has the same length

Indexing data frames

[] can work for indexing data frames, just like matrices:

```
1 example_df <- data.frame(x = c(1, 2, 3),  
2                           y = c('a', 'b', 'c'))  
3 example_df
```

```
  x y  
1 1 a  
2 2 b  
3 3 c
```

```
1 example_df[2, 1]
```

```
[1] 2
```

Indexing data frames

Like lists, `[[]]` and `$` can also be used:

```
1 example_df <- data.frame(x = c(1, 2, 3),  
2                           y = c('a', 'b', 'c'))  
3 example_df
```

```
  x y  
1 1 a  
2 2 b  
3 3 c
```

```
1 example_df$x
```

```
[1] 1 2 3
```

```
1 example_df[["x"]]
```

```
[1] 1 2 3
```

What do you do with a data frame?

Data manipulation

```
1 glimpse(starwars)
```

```
Rows: 87
Columns: 14
$ name      <chr> "Luke Skywalker", "C-3PO", "R2-D2", "Darth Vader",
"Leia Or...
$ height    <int> 172, 167, 96, 202, 150, 178, 165, 97, 183, 182, 188,
180, 2...
$ mass      <dbl> 77.0, 75.0, 32.0, 136.0, 49.0, 120.0, 75.0, 32.0,
84.0, 77...
$ hair_color <chr> "blond", NA, NA, "none", "brown", "brown, grey",
"brown", N...
$ skin_color <chr> "fair", "gold", "white, blue", "white", "light",
"light", "...
$ eye_color  <chr> "blue", "yellow", "red", "yellow", "brown", "blue",
"blue",...
```

What manipulation might I want to do with the `starwars` data?

dp_lyr: Tools for data wrangling



- part of the tidyverse
- provides a “grammar of data manipulation”: useful verbs (functions) for manipulating data
- we will cover the key dp_lyr functions

Verbs for data wrangling

- `select()`: take a subset of the columns (i.e., features, variables)
- `filter()`: take a subset of the rows (i.e., observations)
- `mutate()`: add or modify existing columns
- `arrange()`: sort the rows
- `summarize()`: aggregate the data across rows (e.g., group it according to some criteria)

Creating a subset of the rows

Question: Suppose I only want the droids in the `starwars` data. How would I choose only those rows?

Creating a subset of the rows

Question: Suppose I only want the droids in the `starwars` data. How would I choose only those rows?

```
1 filter(starwars, species == "Droid")
```

```
# A tibble: 6 × 14
```

	name	height	mass	hair_color	skin_color	eye_color	birth_year	sex
	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>
1	C-3PO	167	75	<NA>	gold	yellow	112	none
2	R2-D2	96	32	<NA>	white, blue	red	33	none
3	R5-D4	97	32	<NA>	white, red	red	NA	none
4	IG-88	200	140	none	metal	red	15	none
5	R4-P17	96	NA	none	silver, red	red, blue	NA	none

Creating a subset of the rows

```
1 starwars |>
2   filter(species == "Droid")
```

```
# A tibble: 2 × 14
  name    height  mass hair_color skin_color eye_color birth_year sex
  <chr>    <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr>
1 C-3PO    167    75 <NA>      gold        yellow      112 none
  masculine
2 R2-D2     96    32 <NA>      white, blue red        33 none
  masculine
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

- `|>` is called the *pipe*. It means “take <this>, THEN do <that>”
- `filter` keeps only the rows which satisfy a specific condition

Calculating summary statistics

Question: What is the average height for droids in the dataset?

Calculating summary statistics

Question: What is the average height for droids in the dataset?

```
1 starwars |>
2   filter(species == "Droid") |>
3   summarize(mean_height = mean(height))
```

```
# A tibble: 1 × 1
  mean_height
    <dbl>
1         NA
```

- pipes (|>) can be chained together
- `summarize` calculates summary statistics
- Why am I getting NA?

Handling missing values

```
# A tibble: 6 × 14
  name      height  mass hair_color skin_color eye_color birth_year sex
gender
  <chr>    <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr>
<chr>
1 C-3PO      167    75 <NA>      gold        yellow      112 none
masculi...
2 R2-D2       96    32 <NA>      white, blue red        33 none
masculi...
3 R5-D4       97    32 <NA>      white, red  red        NA none
masculi...
4 IG-88      200   140 none      metal        red        15 none
masculi...
5 R4-P17      96    NA none      silver, red red, blue   NA none
```

```
1 starwars |>
2   filter(species == "Droid") |>
3   summarize(mean_height = mean(height, na.rm=T))
```

```
# A tibble: 1 × 1
  mean_height
  <dbl>
1      131.
```

Calculating summary statistics

Question: What if I want the average height for *humans*?

```
1 starwars |>
2   filter(species == "Droid") |>
3   summarize(mean_height = mean(height, na.rm=T))
```


Calculating summary statistics

Question: What if I want the average height for *humans*?

```
1 starwars |>
2   filter(species == "Human") |>
3   summarize(mean_height = mean(height, na.rm=T))
```

```
# A tibble: 1 × 1
  mean_height
    <dbl>
1       177.
```

Calculating summary statistics

Question: What is the average body mass for *each* species?

Calculating summary statistics

Question: What is the average body mass for *each* species?

```
1 starwars |>
2   group_by(species) |>
3   summarize(mean_height = mean(height, na.rm=T))
```

```
# A tibble: 38 × 2
  species    mean_height
  <chr>         <dbl>
1 Aleena           79
2 Besalisk        198
3 Cerean          198
4 Chagrian        196
5 Clawdite        168
6 Droid          131.
7 Dug            112
8 Ewok             88
9 Geonosian       183
10 Gungan        209.
# i 28 more rows
```

Creating new variables

Question: What is the distribution of the ratio of body mass to height?

Creating new variables

Question: What is the distribution of the ratio of body mass to height?

```
1 starwars |>  
2   mutate(body_ratio = mass/height)
```

Creating new variables

```
1 starwars |>
2   mutate(body_ratio = mass/height) |>
3   group_by(species) |>
4   summarize(mean_ratio = mean(body_ratio, na.rm=T),
5             sd_ratio = sd(body_ratio, na.rm=T))
```

A tibble: 38 × 3

	species	mean_ratio	sd_ratio
	<chr>	<dbl>	<dbl>
1	Aleena	0.190	NA
2	Besalisk	0.515	NA
3	Cerean	0.414	NA
4	Chagrian	NaN	NA
5	Clawdite	0.327	NA
6	Droid	0.453	0.174
7	Dug	0.357	NA
8	Ewok	0.227	NA
9	Geonosian	0.437	NA
10	Gungan	0.351	0.0207

i 28 more rows

Creating new variables

```
1 starwars |>
2   mutate(body_ratio = mass/height) |>
3   group_by(species) |>
4   summarize(mean_ratio = mean(body_ratio, na.rm=T),
5             sd_ratio = sd(body_ratio, na.rm=T),
6             N = n())
```

A tibble: 38 × 4

	species	mean_ratio	sd_ratio	N
	<chr>	<dbl>	<dbl>	<int>
1	Aleena	0.190	NA	1
2	Besalisk	0.515	NA	1
3	Cerean	0.414	NA	1
4	Chagrian	NaN	NA	1
5	Clawdite	0.327	NA	1
6	Droid	0.453	0.174	6
7	Dug	0.357	NA	1
8	Ewok	0.227	NA	1
9	Geonosian	0.437	NA	1
10	Gungan	0.351	0.0207	3

i 28 more rows

Summary so far

- `filter`: choose certain rows
- `summarize`: calculate summary statistics
- `group_by`: group rows together
- `mutate`: create new columns

Class activity

https://sta279-s24.github.io/class_activities/ca_lecture_9.html

