

Lecture 26: C++ and Rcpp

A snippet of C++ code in R

load C++ code into R

```
1 Rcpp::cppFunction('int add(int x, int y, int z) {  
2   int sum = x + y + z;  
3   return sum;  
4 }')  
5  
6 add(1, 2, 3) ← calling the function in R
```

function

declare input types
(in this case, int)

return type

return
the
sum of
3
integers

[1] 6

What is this code doing?

Adding 3 integers together

Function: takes in 3 ints

returns an int

when we make a new object (e.g., sum)
we need to specify its type

Here: we returning "sum"
if sum was not an int, we would get an error

C++ code

```
1 int add(int x, int y, int z) {  
2   int sum = x + y + z;  
3   return sum;  
4 }
```

assignment

What are some differences between C++ and R code?

· C++ : need to specify the type of everything

· naming function,

R:
add <- function(x, y, z) {

}

C++
int add(int x, ...) {

}

· C++ : semicolon to end lines

C++ code

Here's another function:

```
1 int signC(int x) {  
2     if (x > 0) {  
3         return 1;  
4     } else if (x == 0) {  
5         return 0;  
6     } else {  
7         return -1;  
8     }  
9 }
```

Handwritten annotations:

- ← returns an integer (pointing to the `int` return type)
- input: an integer (pointing to the `int x` parameter)
- returning the sign of an integer (pointing to the `if` and `else if` blocks)

What similarities do you notice between C++ and R?

- `if ... else if ... else` is the same in C++ and R
- `==` to test for equality

C++ code

vector object

(behaves similarly to vectors in R)

1 **double** sumC(NumericVector x) {

2 int n = x.size(); ← length of x

3 double total = 0;

4 for(int i = 0; i < n; **++i**) { ← add 1 to i

5 total **+=** x[i]; ← updating total with x[i]

6 }

7 return total; total = total + x[i]

8 }

x.size() gives length of a Numeric Vector

What is this code doing?

Sums up all entries in a vector

for loop: Start at $i = 0$

Stop at $i = n - 1$

increment by 1 each time

$i = 0$, then $i = 1$, then $i = 2$, ...

(In R: $\text{for}(i \text{ in } 1:n) \{ \dots \}$)

In C++, indices start at 0
• $x[i]$ is entry in x


return type
double
eg,
(1.0,
1.5,
1.23, 7...)

Comparing R and C++ speed

```
1 Rcpp::cppFunction('double sumC(NumericVector x) {  
2   int n = x.size();  
3   double total = 0;  
4   for(int i = 0; i < n; ++i) {  
5     total += x[i];  
6   }  
7   return total;  
8 }')  
9  
10 x <- rnorm(1000)  
11 bench::mark(  
12   sum(x), ← Sum function in R  
13   sumC(x) ← or C++ version  
14 )
```

A tibble: 2 × 6

	expression	min	median	`itr/sec`	mem_alloc	`gc/sec`
	<bch:expr>	<bch:tm>	<bch:tm>	<dbl>	<bch:byt>	<dbl>
1	sum(x)	111.75µs	112.12µs	8849.	0B	0
2	sumC(x)	2.33µs	3.25µs	310597.	2.49KB	0

 about 30 times faster

C++ code

```
1 NumericVector col_meansC(NumericMatrix x) {  
2   int n_cols = x.ncol();  
3   int n_rows = x.nrow();  
4   NumericVector col_means(n_cols);  
5  
6   double total = 0;  
7  
8   for(int j = 0; j < n_cols; ++j){  
9     total = 0;  
10    for(int i = 0; i < n_rows; ++i){  
11      total += x(i,j);  
12    }  
13    col_means[j] = total/n_rows;  
14  }  
15  
16  return col_means;  
17 }
```

return a vector
(one entry, i.e. one column
mean, per column)

input is a matrix

getting dimensions of our matrix
(# cols, # rows)

create a vector (of 0s)
variable to store the total of length n_cols

iterating over columns

reset total for each new column

iterating over rows

$x(i,j)$ value at row i , column j

getting the mean
(dividing the sum by # of rows)

Comparing R and C++ speed

```
1 x <- matrix(rnorm(1000*150), ncol=150)
2
3 bench::mark(
4   colMeans(x),
5   col_meansC(x)
6 )
```

A tibble: 2 × 6

	expression	min	median	`itr/sec`	mem_alloc	`gc/sec`
	<bch:expr>	<bch:tm>	<bch:tm>	<dbl>	<bch:byt>	<dbl>
1	colMeans(x)	4.04ms	4.07ms	245.	25.45KB	0
2	col_meansC(x)	123.33µs	127.38µs	7779.	3.71KB	0

↗
quite a bit
faster with
C++ implementation

Some key points

- C++ *always* needs to know the **type** of an object
 - This is true for inputs, outputs, *and* any variables you create
- In C++, indexing begins at 0
- C++ needs a ; at the end of each line
- `NumericVector` objects are the equivalent of vectors in R
- `NumericMatrix` objects are the equivalent of matrices in R

Some useful C++ code

- `.size()` returns the length of a `NumericVector`
- `.ncol()` and `.nrow()` give the numbers of rows and columns for a `NumericMatrix`
- `a += b` is shorthand for `a = a + b`
- `pow()` is used for exponentiation (e.g., `pow(3, 2)`)
- `int` is an integer value, `double` is a decimal value

Class activity

https://sta279-s24.github.io/class_activities/ca_lecture_26.html

